



METIS

Train / Test Splits and Cross Validation

LEARNING GOALS



1. Understand the importance of estimating model generalization error (testing) and practical methodologies for doing so
2. Learn validation / cross-validation strategies for model selection
3. Recognize the difference between validation and testing

Testing: Estimating Generalization Error

METIS



Generalization Error: How well can we expect a model to perform on new data from the same distribution as the training data?

- Predictive models are only *useful* if they can give us good target approximations for samples that we haven't seen before
- Example: Zillow predicts the market value of a home before it's listed for sale, training a model on known listing prices
- So when evaluating models, we should attempt to measure how well they *generalize*, i.e. estimate performance on samples we didn't train on. We call this **testing**.

TESTING, IN PRACTICE



1. Fit model to training data



Train (80%)

Test (20%)

TESTING, IN PRACTICE; cont.



1. Fit model to training data
2. Score model on testing data




Train (80%)

Test (20%)

TEST USUALLY UNDERPERFORMS TRAIN



Model is optimized to perform as well as possible on train, so it's no surprise that it tends to have a worse evaluation score on test (though this is not guaranteed).



→ R^2 : .49




→ R^2 : .43


THE RISK OF OVERFITTING



Out of sample evaluations let us check for overfitting: more complex models can get arbitrarily better at predicting the train data, but will start to fit to spurious patterns and generalize more poorly



→ Degree 4 R^2 : .54



→ Degree 4 R^2 : .42

Validation: Optimizing Our Modeling Choices

METIS

VALIDATION: OPTIMIZING CHOICES



When we construct predictive models, we typically have **many choices**:


- Features: which data columns do we include/exclude or engineer?
- Preprocessing: how should we handle nulls? Should we standardized the scale of the features?
- Hyper-parameters: What degree polynomial regression should we fit? What regularization strength should we use? How does a random forest model compare to a linear regression model?

VALIDATION IN PRACTICE



1. Train candidate models

- Linear Regression
- Polynomial Regression
- Ridge Regression



VALIDATION IN PRACTICE; cont.



1. Train candidate models

- Linear Regression
- Polynomial Regression
- Ridge Regression

2. Score candidates

→ R^2 : .35

→ R^2 : .50

→ R^2 : .25

Train (60%)

Validation (20%)

Test (20%)



VALIDATION IN PRACTICE; cont.



3. Retrain best candidate on train + validation

→ Polynomial Regression




Train + Valid (80%)

Test (20%)

VALIDATION IN PRACTICE; cont.



4. Score final model on test



VALIDATION: KEY CONSIDERATIONS



Validation is not testing: This is a very common pitfall. Once you've used a data set to influence your model choices through direct feedback, it can't be used to give an unbiased estimate of generalization error

Fair comparisons: Candidate models should be compared against the same validation scheme

Data efficiency: after we use a portion of our data for validation, we should reuse it as training data to improve the final model

Cross Validation: Optimizing Our Optimization Of Choices

METIS

CROSS VALIDATION: ADDING MORE RIGOR



Cross-validation is about reliability and efficiency: what if we overfit to an unlucky validation set? Can we use more of our data than just one hold-out for validation?

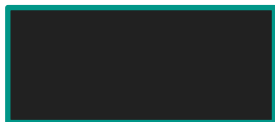
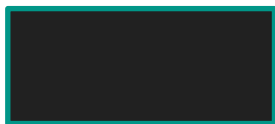

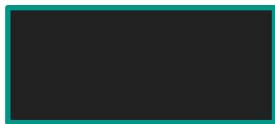
K-Fold partitioning: randomly divide our non-test data into K equal-sized groups. Each group will be used as a validation set once, and we'll compare candidate models via mean scores across all validation scores.

K is usually 5 or 10: depends on problem and size of data, but these are common choices

CROSS VALIDATION IN PRACTICE



Train data:
divided into
5 folds



Do the following,
for each
candidate model

Test, held out

CROSS VALIDATION IN PRACTICE



Train data:
divided into
5 folds



Train on
black folds


Validate on
cyan fold

Validation R^2 scores: .45,

CROSS VALIDATION IN PRACTICE



Train data:
divided into
5 folds



Train on
black folds


Validate on
cyan fold

Validation R^2 scores: .45, .39,

CROSS VALIDATION IN PRACTICE

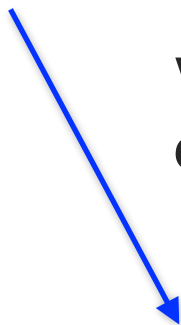


Train data:
divided into
5 folds



Train on
black folds

Validate on
cyan fold

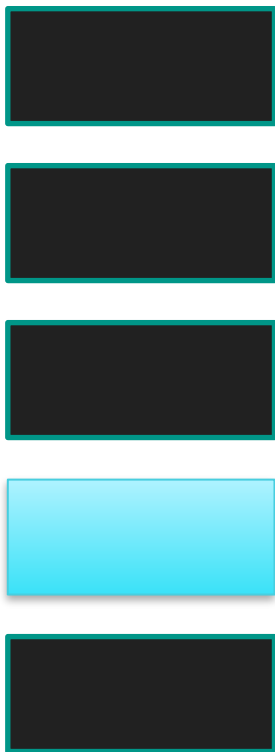


Validation R^2 scores: .45, .39, .57,

CROSS VALIDATION IN PRACTICE



Train data:
divided into
5 folds



Train on
black folds

Validate on
cyan fold

Validation R^2 scores: .45, .39, .57, .49,

CROSS VALIDATION IN PRACTICE




Train data:
divided into
5 folds



Train on
black folds

Validate on
cyan fold







Validation R^2 scores: .45, .39, .57, .49, .50

CROSS VALIDATION IN PRACTICE



Train data:
divided into
5 folds



Produces a set of results for each
candidate model

Linear regression



Validation R^2 scores: .45, .39, .57, .49, .50

Poly regression







Validation R^2 scores: .53, .43, .67, .55, .51

CROSS VALIDATION IN PRACTICE



Train data:
divided into
5 folds



Summarize candidates by mean
score, select best

Linear regression



5-Fold validation mean R^2 score: .48

Poly regression




5-Fold validation mean R^2 score: .54

CROSS VALIDATION IN PRACTICE



Train data:
recombined



Polynomial regression
selected as best
candidate model

**Poly regression,
retrained on all data,
final score on test**



Test, held out

R^2 : .48

Validation And Testing: Recap


METIS

WORKFLOW METHOD 0: Train/Test



This is what you'll see in a lot of teaching examples to demonstrate how an algorithm works, but is not good to use in practice.

Instead, use one of the following methods.




WORKFLOW METHOD 1: Train/Valid/Test



Collect set of candidate models. Fit each on train, score on validation, select final model via best validation score

Retrain final model on train + validation, report score on test as estimate of generalization error



WORKFLOW METHOD 2: CV/Test



Collect set of candidate models. Run each through a K-fold CV loop, select final model via best mean validation score

Retrain final model on combined folds, report score on test as estimate of generalization error



VALIDATION VS. CV - WHEN TO USE?



Simple validation is significantly faster and often representative enough when working with very large samples (~millions+)

Cross validation is more appropriate with small-medium size data or when variance in results between different validation sets is high

SUMMARY



Training

In sample

Model building

Optimize model
parameters (fit)

Validation

Out of sample

Feedback to model
selection

Optimize choices:
features and model
hyper-parameters

Testing

Out of sample

No feedback to model
selection

Final estimate of model
generalization error

VALIDATION SO WOW

SUCH MODEL SELECTION

**VERY NOT ESTIMATE
OF GENERALIZATION ERROR**